

Table of Contents

Matplotlib module

<https://matplotlib.org/stable/index.html> (<https://matplotlib.org/stable/index.html>)







Installation Documentation Examples Tutorials Contributing

[home](#) | [contents](#) » [Matplotlib: Python plotting](#)

Matplotlib: Visualization with Python

Matplotlib is a comprehensive library for creating static, animated, and interactive visualizations in Python.



Matplotlib makes easy things easy and hard things possible.

Create

- Develop **publication quality plots** with just a few lines of code
- Use **interactive figures** that can zoom, pan, update...

Customize

- **Take full control** of line styles, font properties, axes properties...
- **Export and embed** to a number of file formats and interactive environments

Extend

- Explore tailored functionality provided by **third party packages**
- Learn more about Matplotlib through the many **external learning resources**

Check the gallery for getting a glimpse of what you can do with matplotlib!

<https://matplotlib.org/stable/gallery/index.html>
(<https://matplotlib.org/stable/gallery/index.html>)

Need to **import the module** and define some convenient abbreviations

In [1]:

```
import matplotlib.pyplot as plt # import the module

import matplotlib as mpl
mpl.rcParams['figure.dpi']= 130 # set the resolution to x dpi
```

Let's read a **CSV file** with financial data. We want to inspect the behaviors of a few indicators, such as the closing value, the opening value, and the trade volume.

In [2]:

```
f = open('AAPL.csv')
```

Let's deal with the file *without* using the methods from the `csv` module (just for the sake of playing a bit with string methods!)

The first line of the file, the header, contains the names/descriptions of the fields (the columns of the dataset)

In [3]:

```
header = f.readline()
```

In [4]:

```
header
```

Out[4]:

```
'Date,Open,High,Low,Close,Adj Close,Volume\n'
```

It's a *comma* separated values file, therefore, we can extract the individual fields by splitting on `' , '`

In [5]:

```
fieldnames = header.split(' , ')
```

In [6]:

```
fieldnames
```

Out[6]:

```
['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume\n']
```

Not exactly what we'd like to have: there's a `\n` character that we don't want.

Moreover, what about the potential presence of white spaces?

In [7]:

```
f.seek(0)
header = f.readline()
fieldnames = header[:-1].split(',')
fieldnames
```

Out[7]:

```
['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume']
```

We can read now the file record by record and store the values of the columns `Date`, `Open`, `Close`, `Volume` in different lists (note that we will exploit the knowledge of what content is reported in a column).

Apart from `Date`, all values (that are read as strings) are numeric, such that are casted as floats.

In [8]:

```
open_val = []
close_val = []
volume = []
dates = []

for record in f:
    fields = record.split(',')

    #print(fields)

    dates.append(fields[0])
    open_val.append( float(fields[1]) )
    close_val.append( float(fields[4]) )
    volume.append( float(fields[6]) )
```

In [9]:

```
#open_val
```

How many records in total?

In [10]:

```
len(close_val)
```

Out[10]:


253

How do we check what has happened for the Apple shares? How do we inspect the data visually? Let's make **data plots**!

In [11]:

```
plt.title("APPLE stock quotes in year 2000")  # the title to give to the figure
plt.xlabel("Date index")  # x-axis label
plt.ylabel("Closing value (in USD)")  # y-axis label
plt.plot(close_val)
plt.show()
```

APPLE stock quotes in year 2000



In [12]:

```
for p in [open_val, close_val, volume]:

    plt.title("APPLE stock quotes in year 2000")  # the title to give to the figure

    plt.xlabel("Date index")  # x-axis label

    if p == open_val:
        ylabel = "Opening value"
        color = 'red'

    elif p == close_val:
        ylabel = "Closing value"
        color = 'blue'


    else:
        ylabel = "Volume"
        color = 'green'

    plt.ylabel(ylabel)  # y-axis label


    plt.plot(p, color = color)

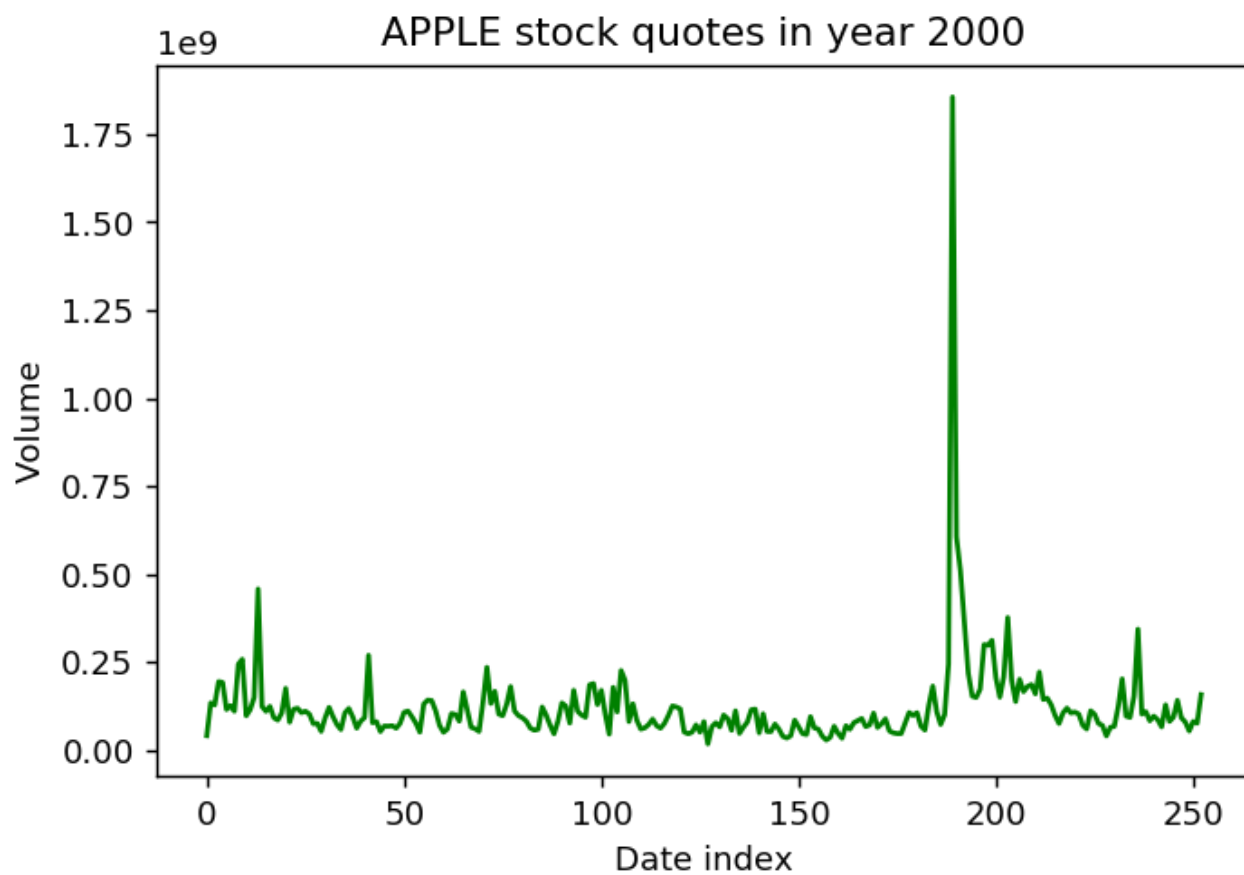
plt.show()
```

APPLE stock quotes in year 2000



APPLE stock quotes in year 2000






Closing and opening values look similar, can we compare them in the same data plot?
\$to\$ **Multiple plots**

In [13]:


```
plt.title("APPLE: Opening and Closing values")
plt.xlabel("Date index")
plt.ylabel("Value (USD)")
plt.plot(close_val)
plt.plot(open_val)
plt.show()
```



Can we improve readability by controlling **style / width of lines** and adding a **legend**?

```
In [14]:
```

```
plt.title("APPLE: Opening vs. Closing values")  
plt.xlabel("Date index")  
plt.ylabel("Value (USD)")  
plt.plot(close_val, linewidth=1, label = 'Closing')  
plt.plot(open_val, linewidth=0.7, linestyle='dashed', label = 'Opening')  
plt.legend()  
plt.show()
```




Different linestyles for the lines:

Solid: `'_'`

Dotted: `'.'`

Dashed: `'--'`

Dashdot: `'_.'`



What about the **dates** on the x axis? Can we write them in a nice way? `\to` **Control ticks/labels on the axes**

```
In [15]:
```

```
plt.title("APPLE: Traded volumes")

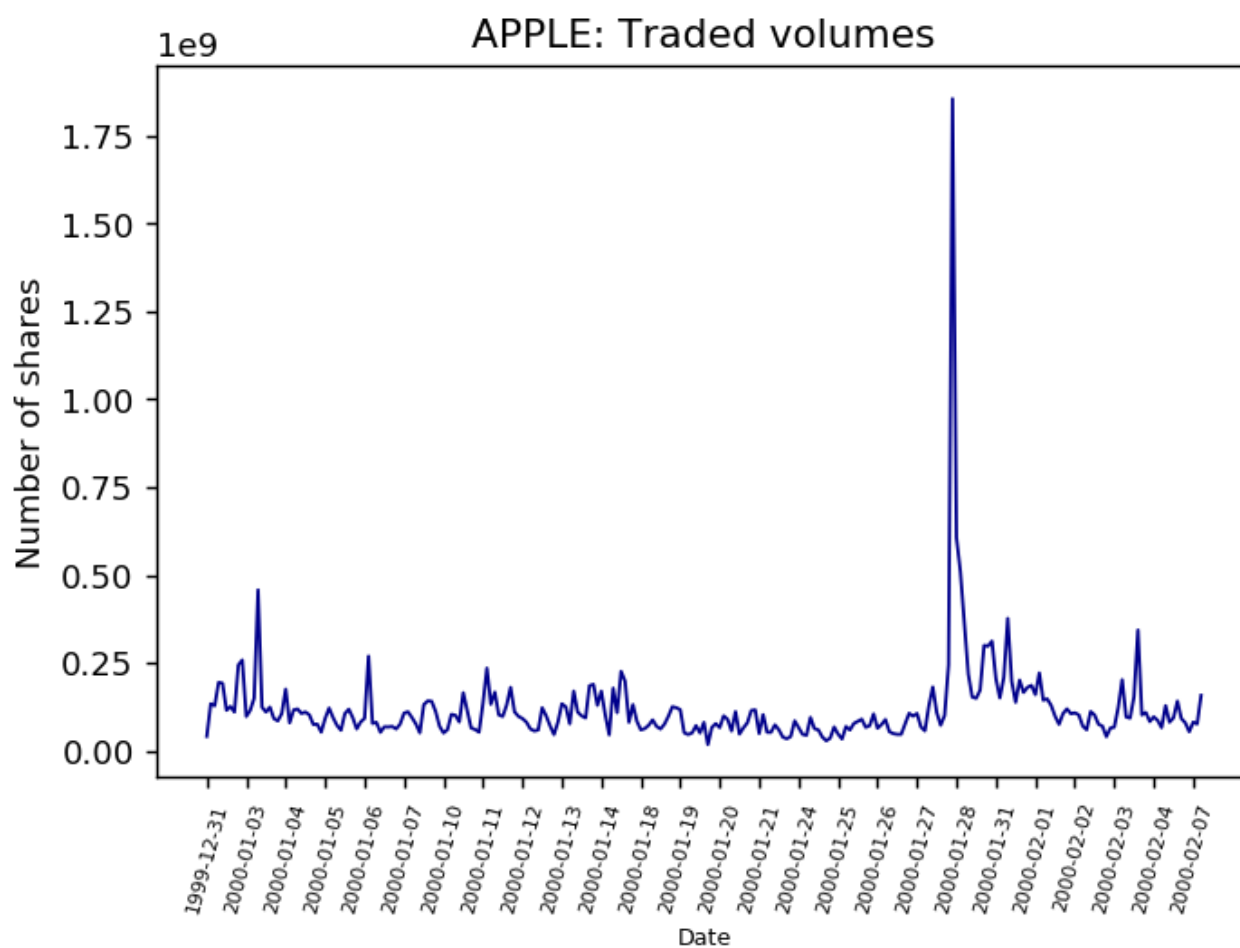
plt.xlabel("Date", fontsize=7)

plt.ylabel("Number of shares")

n_values = len(volume)
plt.xticks(range(0, n_values, 10), dates,
           fontsize=6, rotation=75)

plt.plot(volume, linewidth=1, color = 'darkblue')


plt.show()
```



So far we have only plotted one time series of data. What about correlation between different series of paired data? **\$(x,y)\$ data plots** showing y vs. x **to** **scatter plots**

In [16]:

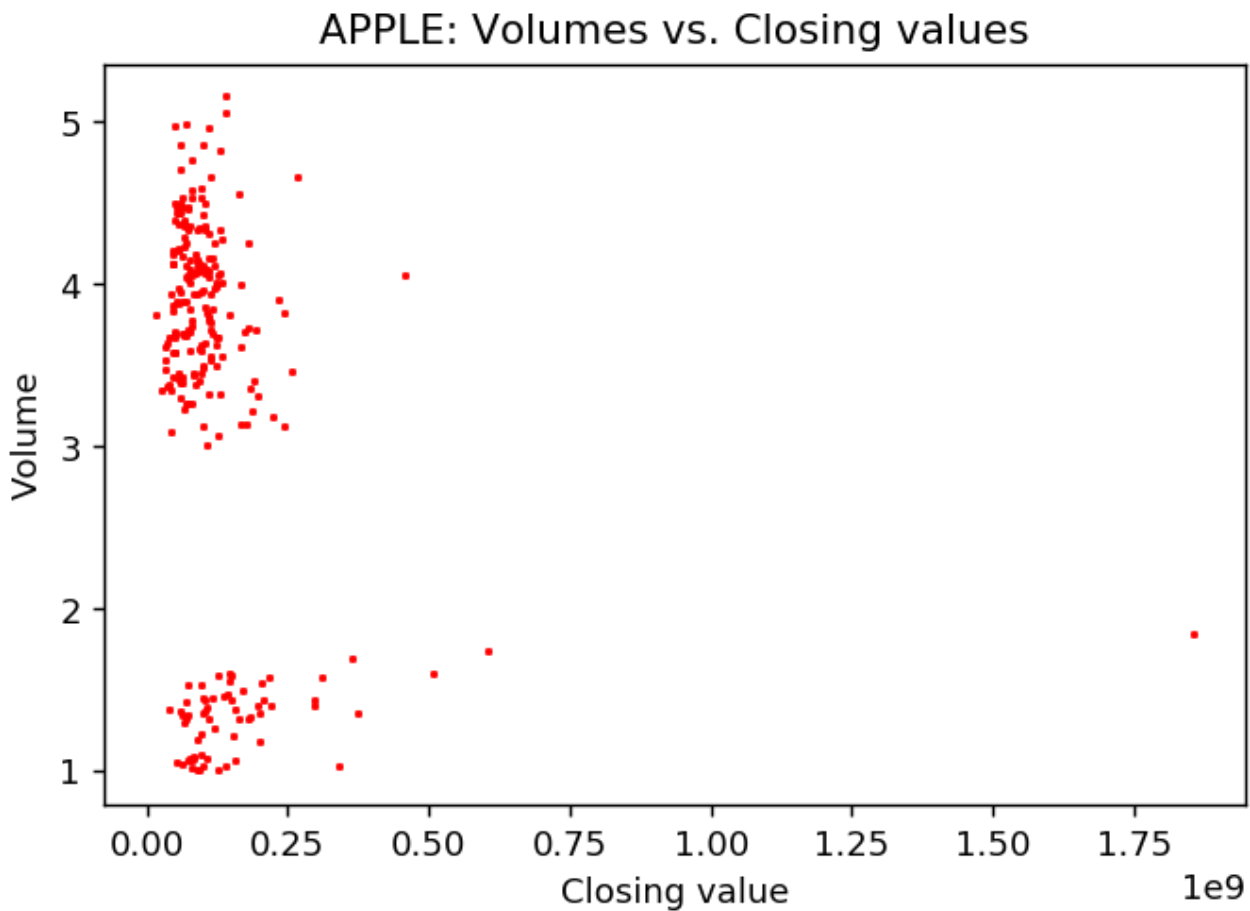
```
plt.title("APPLE: Closing vs. Opening values (a Scatter plot)")
plt.xlabel("Opening value")
plt.ylabel("Closing value")
plt.scatter(open_val, close_val, marker='o', s=1, color='blue')
plt.show()
```



Volumes vs. Closing values, and using a different marker for the points.

In [17]:

```
plt.title("APPLE: Volumes vs. Closing values")
plt.xlabel("Closing value")
plt.ylabel("Volume")
plt.scatter(volume, close_val, marker='s', s=1, color='red')
plt.show()
```



How many **different markers** can we use? A lot! A the named parameter `s` defines the size of the marker. Try them out to see the effect!

In [18]:

```
# Many different markers are available for the shape of the data points
# =====
# character      description
# =====
# ``'.'``       point marker
# ``','``       pixel marker
# ``'o'``       circle marker
# ``'v'``       triangle_down marker
# ``'^'``       triangle_up marker
# ``'<'``       triangle_left marker
# ``'>'``       triangle_right marker
# ``'1'``       tri_down marker
# ``'2'``       tri_up marker
# ``'3'``       tri_left marker
# ``'4'``       tri_right marker
# ``'s'``       square marker
# ``'p'``       pentagon marker
# ``'*'``       star marker
# ``'h'``       hexagon1 marker
# ``'H'``       hexagon2 marker
# ``'+'``       plus marker
# ``'x'``       x marker
# ``'D'``       diamond marker
# ``'d'``       thin_diamond marker
# ``'|'``       vline marker
# ``'_'``       hline marker
# =====
```

Plots with **lines and point markers?**

```
In [19]:
```

```
plt.title("APPLE: Traded volumes")

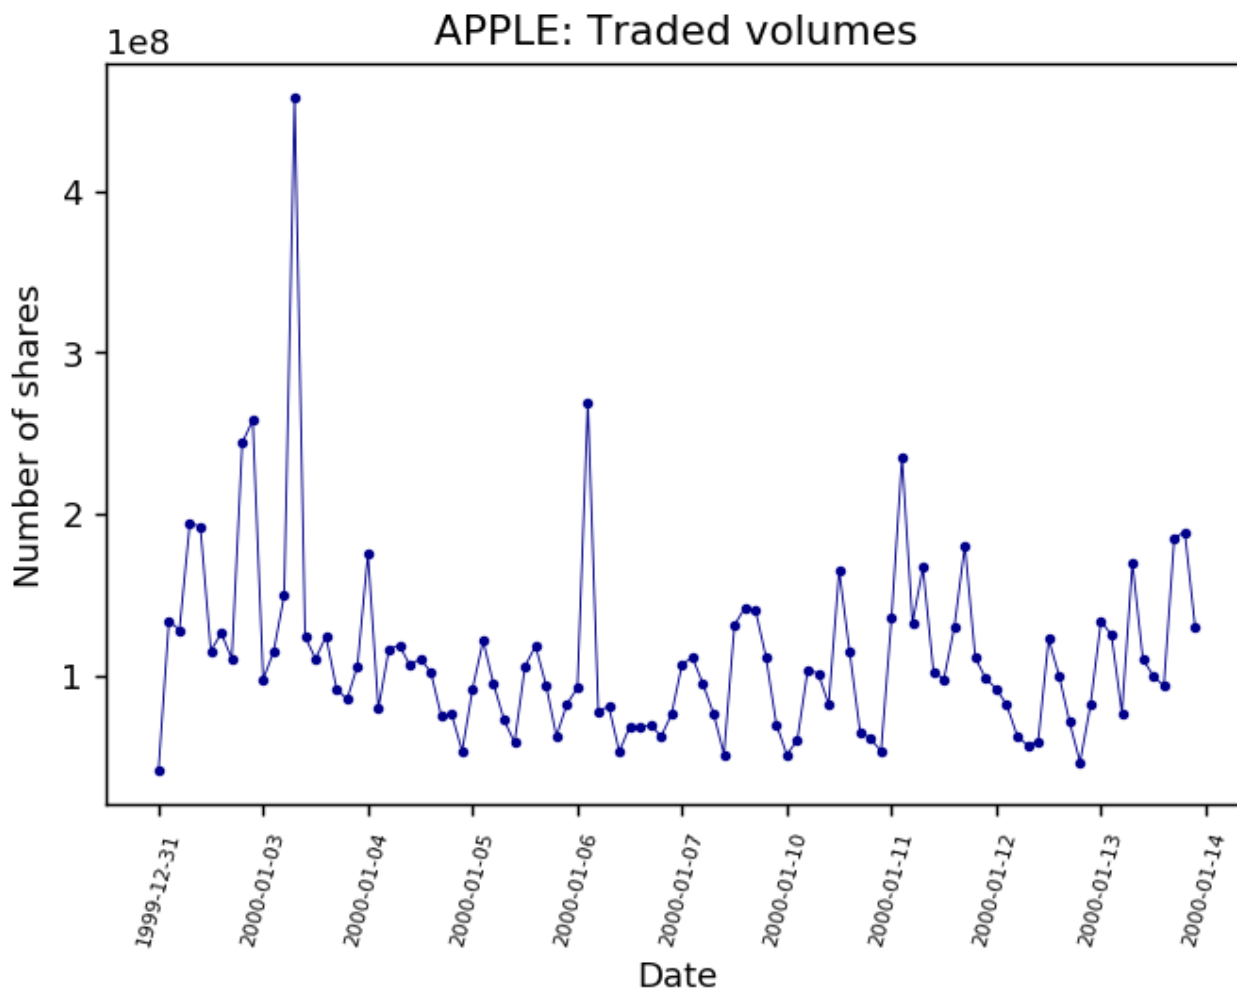
plt.xlabel("Date")

plt.ylabel("Number of shares")

n_values = len(volume)
plt.xticks(range(0, n_values, 10), dates,
           fontsize=6, rotation=75)

plt.plot(volume[0:100],
         marker = 'o', markersize=2,
         linewidth=0.5, color = 'darkblue')

plt.show()
```



We can also change the **colors of the markers (borders and inside)**.

In [20]:

```
plt.title("APPLE: Traded volumes")


plt.xlabel("Date")

plt.ylabel("Number of shares")

n_values = len(volume)
plt.xticks(range(0, n_values, 10), dates,
           fontsize=6, rotation=75)

plt.plot(volume[0:100],
         marker = 'o', markersize=3,
         markerfacecolor='r', markeredgecolor='blue',
         linewidth=0.5, color = 'gray')

plt.show()
```



Can we add a **grid** to read the values more clearly?

```
In [21]:
```

```
plt.title("APPLE: Traded volumes")

plt.xlabel("Date")


plt.ylabel("Number of shares")

n_values = len(volume)
plt.xticks(range(0, n_values, 10), dates,
           fontsize=6, rotation=75)

plt.plot(volume[0:100],
         marker = 'o', markersize=1,
         linewidth=0.5, color = 'red')

plt.grid()

plt.show()
```



What about **resizing** the plots and changing the font size of labels, etc.?

In [22]:

```
plt.figure ( figsize = (3,2))

scaled_fontsize = 7

plt.title("APPLE: Closing values", fontsize=scaled_fontsize)

plt.xlabel("Date", fontsize=scaled_fontsize)

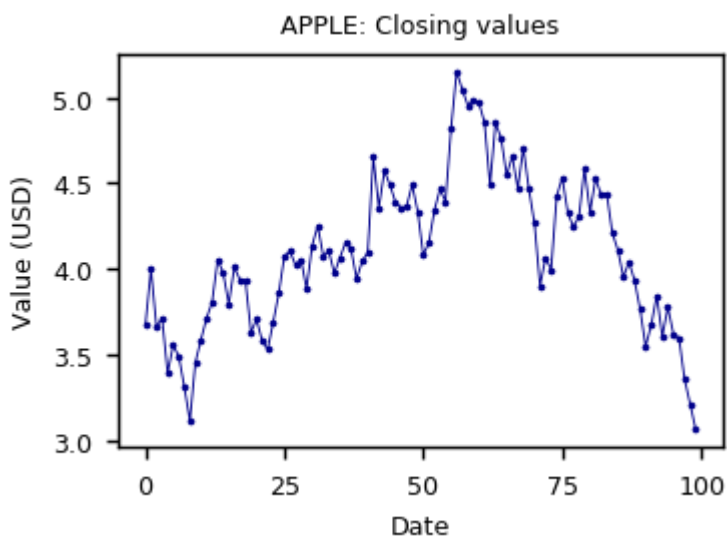
plt.ylabel("Value (USD)", fontsize=scaled_fontsize)

plt.plot(close_val[0:100],
         marker = 'o', markersize=1,
         linewidth=0.5, color = 'darkblue')

plt.xticks(fontsize=scaled_fontsize)

plt.yticks(fontsize=scaled_fontsize)

plt.show()
```



The cells below haven't been presented during the class!

They are there to let you give it a look (for the project), in any case we'll go through them during the next lecture.

Let's read another dataset, reporting **data about mall customers**, and let's use this time `csv` methods (it will be easier than before!)

In [23]:

```
import csv

f = open('Mall_Customers.csv')

f_csv = csv.reader(f)
```

In [24]:

```
fieldnames = next(f_csv)
print(fieldnames)

print(len(fieldnames))

['CustomerID', 'Gender', 'Age', 'Annual Income (k$)', 'Spending Score (1-100)']
5
```

In [25]:

```
gender = []
age = []
income = []
score = []
customer_id = []

for record in f_csv:
    customer_id.append(record[0])

    gender.append(record[1].strip()) # .strip() removes all extra white spaces!

    age.append( int(record[2]) )
    income.append( int(record[3]) )
    score.append( int(record[4]) )
```

In [26]:

```
#age
```

In [27]:


```
#gender
```

This is **not a time series** of data!

What about the distribution of the values of income?

In [28]:

```
plt.title("Mall customers: Income")
plt.xlabel("Customer id")
plt.ylabel("Income")
plt.scatter(customer_id, income, marker='o', s=0.5, color='red')
plt.xticks(customer_id[1::10], fontsize=6)
plt.show()
```




It looks like the income grows with the customer id! (values have been sorted vs. the income)

What about the shopping scores?

In [29]:


```
plt.title("Mall customers: Scores")
plt.xlabel("Customer id")
plt.ylabel("Score")
plt.scatter(customer_id, score, marker='o', s=0.5, color='blue')
plt.xticks(customer_id[1::10], fontsize=6)
plt.show()
```



Aggregating the data to check what the **distribution** of the scores look like in the customer population can be useful \$to\$ **Histogram**

In [30]:


```
plt.title("Mall customers: Histogram Distribution of Scores")
plt.xlabel("Score")
plt.ylabel("Occurrences")
plt.hist(score, bins=10, color='blue')
#plt.xticks(customer_id[1::10], fontsize=6)
plt.show()
```



What about the distribution of the age of the customers?

In [31]:

```
plt.title("Mall customers: Histogram Distribution of Customer Age")
plt.xlabel("Age")
plt.ylabel("Occurrences")
plt.hist(age, bins=15, align='left',
         color='orange')
plt.xticks(range(15, 75, 5), fontsize=8)
plt.show()
```



What about the distribution of male vs. female?

In this case we have two values: number of male and number of female customers.

We need to count them from the `gender` list!

In [32]:

```
n_male = gender.count('Male')  
n_female = gender.count('Female')  
  
n_male, n_female
```

Out[32]:

(88, 112)

How do we effectively show these proportions? (a histogram is not really appropriate for showing proportions)

- **Pie chart**
- **Bar chart**


In [33]:

```
plt.title("Proportions of Male and Female customers")

plt.pie([n_male, n_female],
        labels = ['Male', 'Female'],
        autopct="%.1f%%") # this says to use one decimal digit and to use percentages

plt.show()
```

Proportions of Male and Female customers




Some beautifying of the pie chart.

`shadow` adds a shadowing, while `explode` is tuple, where if the *i*-th value is greater than 0, it indicates that the *i*-th slice of the pie will be detached from the rest of the pie (at a distance proportional to the indicated value).

In [34]:

```
plt.title("Proportions of Male and Female customers")  
  
plt.pie([n_male, n_female],  
        labels = ['Male', 'Female'], shadow=True, explode = (0, 0.1),  
        autopct="%.1f%%")  
  
plt.show()
```

Proportions of Male and Female customers



Can we make a pie chart for the customer age?

Yes, but we need to define age ranges, group / count the data accordingly, and store/define explanatory labels.

Let's consider 10-year ranges, starting from the minimal age up to the maximal age in the dataset.

We'll use a dictionary data structure.

In [35]:

```
import math

min_age = min(age)
max_age = max(age)

age_interval = 10

age_ranges = (max_age - min_age) / age_interval

age_ranges = math.ceil(age_ranges)
print(min_age, max_age, age_ranges)
```

18 70 6

In [36]:

```
age_dict = {}

for r in range(age_ranges):
    range_min = min_age + r * age_interval
    range_max = range_min + age_interval - 1
    range_str = str(range_min) + '-' + str(range_max)
    age_dict[ range_str ] = [range_min, range_max, 0]
```

In [37]:

```
age_dict
```

Out[37]:

```
{'18-27': [18, 27, 0],
 '28-37': [28, 37, 0],
 '38-47': [38, 47, 0],
 '48-57': [48, 57, 0],
 '58-67': [58, 67, 0],
 '68-77': [68, 77, 0]}
```

In [38]:

```
for v in age:
    for r in age_dict:
        if v >= age_dict[r][0] and v <= age_dict[r][1]:
            age_dict[r][2] += 1
            break
```

In [39]:

```
age_dict
```

Out[39]:

```
{'18-27': [18, 27, 46],  
 '28-37': [28, 37, 61],  
 '38-47': [38, 47, 36],  
 '48-57': [48, 57, 31],  
 '58-67': [58, 67, 20],  
 '68-77': [68, 77, 6]}
```

Now we can show the counts in the selected ranges of age as proportions using a pie chart.

Counts are stored in the list `age_counts`

In [40]:

```
plt.title("Proportions of different Age ranges in mall customers")

age_counts = []
for r in age_dict.values():
    age_counts.append(r[2])


# this is to identify the slice with the largest proportion to explode it
max_range = -1
max_range_idx = -1
for i,r in enumerate(age_dict.values()):
    if r[2] > max_range:
        max_range = r[2]
        max_range_idx = i

explode_flag = [0] * len(age_dict)
explode_flag[max_range_idx] = 0.1

plt.pie(age_counts,
        labels = list(age_dict.keys()),
        #shadow=True,
        explode = explode_flag,
        autopct="%.1f%%")

plt.show()
```


Proportions of different Age ranges in mall customers




We can display the same data using a **bar chart**.

In [41]:

```
plt.title("Bar Chart for Age ranges")

plt.bar(range(1, len(age_counts)+1),          # position of the bars
        age_counts,                          # value/height of the bars
        tick_label = list(age_dict.keys()))    # what to display at the x ticks

plt.show()
```



In []: